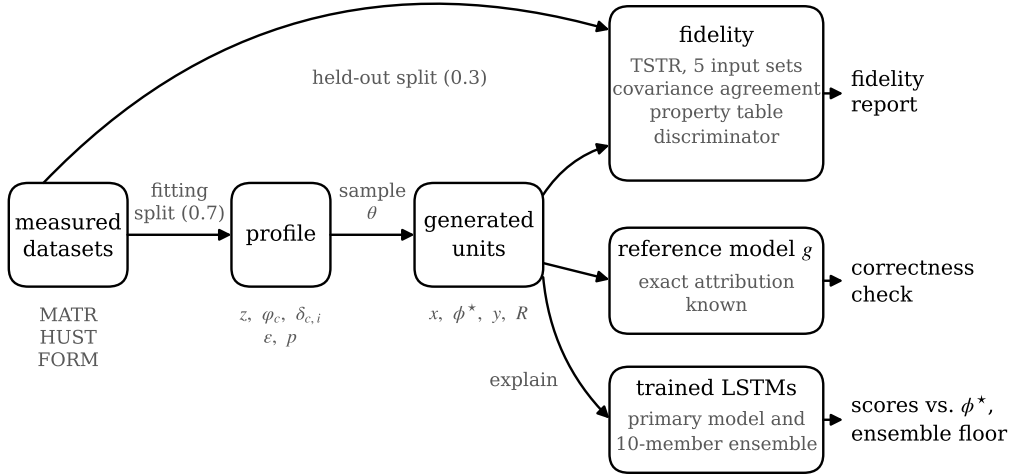
held-out split (0.3)
fidelity
TSTR, 5 input sets
covariance agreement
property table
discriminator
fidelity report
measured datasets
fitting split (0.7)
profile
sample θ
generated units
reference model g
exact attribution known
correctness check
MATR
HUST
FORM
$z,\ \varphi_c,\ \delta_{c,i}$
$\varepsilon,\ p$
$x,\ \phi^\star,\ y,\ R$
explain
trained LSTMs
primary model and 10-member ensemble
scores vs. $\phi^\star$, ensemble floor